

RAG - Retrieval Augmented Generation

RAG combines retrieval and generation to improve LLM answers.

Hybrid Search: Vector + BM25 keyword search

Reranker: Cross-Encoder re-ranks results for precision.

Agentic RAG: LLM decides when and how to retrieve.

Applications: Doc QA, Customer Service, Code Search